

BRACT's

Vishwakarma Institute of Technology, Pune

Practical Implementation Sheet

Department: CSE-AI Semester: 5 Academic Year: 2026-27

Division: C Batch: 3 Practical No: 4

Roll no: 68 PRN No: 12415154 Name of Student: Kalpesh Patil

Subject Name : Deep Learning

Problem Statement : Design and implement a Convolutional Neural Network (CNN) for image classification using the Tomato or Soybean disease dataset.

Algorithm :

- Import required libraries such as TensorFlow, Keras, NumPy, Matplotlib, Seaborn, and Scikit-learn.
- Import the PlantVillage dataset directly from Kaggle using kagglehub.
- Select only the Tomato disease classes from the dataset.
- Select a limited and balanced number of images from each class for faster training.
- Resize images to 96×96 pixels and normalize pixel values.
- Shuffle and divide the dataset into training, validation, and testing sets.
- Create a TensorFlow data pipeline using batching, prefetching, and AUTOTUNE.
- Apply data augmentation using random flipping, rotation, and zooming.
- Build a CNN using convolutional, ReLU, max-pooling, flatten, dense, and dropout layers.
- Compile the model using the Adam optimizer and sparse categorical cross-entropy loss.
- Train the CNN using the training dataset and validate it using the validation dataset.
- Plot training and validation accuracy and loss.
- Evaluate the model on the test dataset.
- Generate a classification report and confusion matrix.
- Display actual and predicted disease classes with prediction confidence.
- Save the trained CNN model.

Code :

```
import os
import kagglehub
import tensorflow as tf
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
```

```
from tensorflow.keras import layers, models
from sklearn.metrics import classification_report, confusion_matrix
```

```
IMG_SIZE = (96, 96)
BATCH_SIZE = 32
SEED = 42
EPOCHS = 8
IMAGES_PER_CLASS = 500
```

```
print("Loading dataset from Kaggle...")

DATASET_PATH = kagglehub.dataset_download(
    "emmarex/plantdisease"
)

DATASET_DIR = os.path.join(
    DATASET_PATH,
    "PlantVillage"
)

print("Dataset loaded.")
```

```
all_classes = sorted([
    folder for folder in os.listdir(DATASET_DIR)
    if os.path.isdir(
        os.path.join(DATASET_DIR, folder)
    )
])

tomato_classes = [
    folder for folder in all_classes
    if folder.startswith("Tomato")
]

print("\nTomato classes:")

for i, name in enumerate(tomato_classes):
    print(i, ":", name)
```

```
X = []
y = []

print("\nLoading images...")

for class_index, class_name in enumerate(tomato_classes):

    class_path = os.path.join(
        DATASET_DIR,
        class_name
    )

    image_files = [
        f for f in os.listdir(class_path)
        if f.lower().endswith(
            (".jpg", ".jpeg", ".png")
        )
    ]

    np.random.seed(SEED)

    np.random.shuffle(image_files)

    image_files = image_files[
        :IMAGES_PER_CLASS
```

```
print(
    class_name,
    "->",
    len(image_files),
    "images"
)

for file in image_files:

    image_path = os.path.join(
        class_path,
        file
    )

    image = tf.keras.utils.load_img(
        image_path,
        target_size=IMG_SIZE
    )

    image = tf.keras.utils.img_to_array(
        image
    )
```

```
X.append(image)
y.append(class_index)

X = np.array(X, dtype=np.float32)
y = np.array(y)

print("\nTotal images:", len(X))
```

```
X = X / 255.0
```

```
rng = np.random.default_rng(SEED)

indices = rng.permutation(len(X))

X = X[indices]
y = y[indices]
```

```
total = len(X)

train_end = int(total * 0.70)
val_end = int(total * 0.85)

X_train = X[:train_end]
y_train = y[:train_end]

X_val = X[train_end:val_end]
y_val = y[train_end:val_end]

X_test = X[val_end:]
y_test = y[val_end:]

print("\nDataset Split:")
print("Training:", len(X_train))
print("Validation:", len(X_val))
print("Testing:", len(X_test))
```

```
train_ds = tf.data.Dataset.from_tensor_slices(
    (X_train, y_train)
)

val_ds = tf.data.Dataset.from_tensor_slices(
    (X_val, y_val)
)

test_ds = tf.data.Dataset.from_tensor_slices(
    (X_test, y_test)
)

AUTOTUNE = tf.data.AUTOTUNE

train_ds = (
    train_ds
    .shuffle(1000, seed=SEED)
    .batch(BATCH_SIZE)
    .prefetch(AUTOTUNE)
)
```

```

val_ds = (
    val_ds
    .batch(BATCH_SIZE)
    .prefetch(AUTOTUNE)
)

test_ds = (
    test_ds
    .batch(BATCH_SIZE)
    .prefetch(AUTOTUNE)
)

```

```

# Visualize Dataset
# =====

plt.figure(figsize=(10, 8))

for images, labels in train_ds.take(1):

    for i in range(9):

        plt.subplot(3, 3, i + 1)

        plt.imshow(images[i])

        plt.title(
            tomato_classes[
                labels[i].numpy()
            ]

            plt.axis("off")

plt.suptitle(
    "Sample Tomato Disease Images"
)

plt.tight_layout()

```

```

# =====
# Data Augmentation
# =====

augmentation = tf.keras.Sequential([
    layers.RandomFlip("horizontal"),
    layers.RandomRotation(0.1),
    layers.RandomZoom(0.1)
])

```

```
# =====
# CNN Model
# =====

model = models.Sequential([

    layers.Input(
        shape=(96, 96, 3)
    ),

    augmentation,

    layers.Conv2D(
        16,
        3,
        activation="relu"
    ),

    layers.MaxPooling2D(),

    layers.Conv2D(
        32,
        3,
        activation="relu"
    ),
```

```
        layers.MaxPooling2D(),

        layers.Conv2D(
            64,
            3,
            activation="relu"
        ),

        layers.MaxPooling2D(),

        layers.Flatten(),

        layers.Dense(
            64,
            activation="relu"
        ),

        layers.Dropout(0.3),

        layers.Dense(
            len(tomato_classes),
            activation="softmax"
        )
    ])

```

```
model.compile(
    optimizer="adam",
    loss="sparse_categorical_crossentropy",
    metrics=["accuracy"]
)
```

```

history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=EPOCHS
)

```

```

plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)

plt.plot(
    history.history["accuracy"],
    label="Training Accuracy"
)

plt.plot(
    history.history["val_accuracy"],
    label="Validation Accuracy"
)

plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.title("Training vs Validation Accuracy")
plt.legend()

plt.subplot(1, 2, 2)

plt.plot(
    history.history["loss"],

```

```

    label="Training Loss"
)

plt.plot(
    history.history["val_loss"],
    label="Validation Loss"
)

plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.title("Training vs Validation Loss")
plt.legend()

plt.tight_layout()
plt.show()

```

```

test_loss, test_accuracy = model.evaluate(
    test_ds,
    verbose=1
)

print("\n=====")
print("MODEL PERFORMANCE")
print("=====")

print(
    "Test Accuracy:",
    round(test_accuracy * 100, 2),
    "%"
)

print(
    "Test Loss:",
    round(test_loss, 4)
)

```

```

y_true = []
y_pred = []

for images, labels in test_ds:

    predictions = model.predict(
        images,
        verbose=0
    )

    y_true.extend(
        labels.numpy()
    )

    y_pred.extend(
        np.argmax(
            predictions,
            axis=1
        )
    )

```

```

print("\nClassification Report:\n")

print(
    classification_report(
        y_true,
        y_pred,
        target_names=tomato_classes,
        zero_division=0
    )
)

```

```

cm = confusion_matrix(
    y_true,
    y_pred
)

plt.figure(figsize=(10, 8))

sns.heatmap(
    cm,
    annot=True,
    fmt="d",
    cmap="Blues",
    xticklabels=tomato_classes,
    yticklabels=tomato_classes
)

plt.xlabel("Predicted Class")
plt.ylabel("Actual Class")
plt.title("Confusion Matrix")

```

```

plt.figure(figsize=(12, 9))

for images, labels in test_ds.take(1):

    predictions = model.predict(
        images,
        verbose=0
    )

    for i in range(
        min(9, len(images))
    ):

        predicted = np.argmax(
            predictions[i]
        )

        confidence = (
            np.max(predictions[i]) * 100
        )

        actual_name = tomato_classes[
            labels[i].numpy()
        ]

        predicted_name = tomato_classes[

```

```

        predicted
    ]

    plt.subplot(3, 3, i + 1)

    plt.imshow(images[i])

    color = (
        "green"
        if labels[i].numpy() == predicted
        else "red"
    )

    plt.title(
        f"Actual: {actual_name}\n"
        f"Predicted: {predicted_name}\n"
        f"Confidence: {confidence:.1f}%",
        color=color,
        fontsize=8
    )

    plt.axis("off")

```

```

plt.suptitle(
    "CNN Prediction Results"
)

plt.tight_layout()
plt.show()

```

```

model.save(
    "tomato_disease_cnn.keras"
)

print("\nModel saved successfully.")

```

Output :

```

Loading dataset from Kaggle...
Using Colab cache for faster access to the 'plantdisease' dataset.
Dataset loaded.

Tomato classes:
0 : Tomato_Bacterial_spot
1 : Tomato_Early_blight
2 : Tomato_Late_blight
3 : Tomato_Leaf_Mold
4 : Tomato_Septoria_leaf_spot
5 : Tomato_Spider_mites_Two_spotted_spider_mite
6 : Tomato__Target_Spot
7 : Tomato__Tomato_YellowLeaf__Curl_Virus
8 : Tomato__Tomato_mosaic_virus
9 : Tomato_healthy

Loading images...
Tomato_Bacterial_spot -> 500 images
Tomato_Early_blight -> 500 images
Tomato_Late_blight -> 500 images
Tomato_Leaf_Mold -> 500 images
Tomato_Septoria_leaf_spot -> 500 images
Tomato_Spider_mites_Two_spotted_spider_mite -> 500 images
Tomato__Target_Spot -> 500 images
Tomato__Tomato_YellowLeaf__Curl_Virus -> 500 images

```



```
Tomato_Septoria_leaf_spot -> 500 images
Tomato_Spider_mites_Two_spotted_spider_mite -> 500 images
Tomato_Target_Spot -> 500 images
Tomato_Tomato_YellowLeaf_Curl_Virus -> 500 images
Tomato_Tomato_mosaic_virus -> 373 images
Tomato_healthy -> 500 images
```

Total images: 4873

Dataset Split:

Training: 3411

Validation: 731

Testing: 731

Sample Tomato Disease Images

Tomato_Bacterial_spot




Tomato_Bacterial_spot



Tomato_Tomato_mosaic_virus



Tomato_Late_blight



Tomato_healthy



Tomato_Spider_mites_Two_spotted_spider_mite



Tomato_Leaf_Mold



Tomato_Early_blight



Tomato_Late_blight



Model: "sequential_1"

Layer (type)	Output Shape	Param #
sequential (Sequential)	(None, 96, 96, 3)	0
conv2d (Conv2D)	(None, 94, 94, 16)	448
max_pooling2d (MaxPooling2D)	(None, 47, 47, 16)	0
conv2d_1 (Conv2D)	(None, 45, 45, 32)	4,640
max_pooling2d_1 (MaxPooling2D)	(None, 22, 22, 32)	0
conv2d_2 (Conv2D)	(None, 20, 20, 64)	18,496
max_pooling2d_2 (MaxPooling2D)	(None, 10, 10, 64)	0
flatten (Flatten)	(None, 6400)	0
dense (Dense)	(None, 64)	409,664
dropout (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 10)	650

Total params: 433,898 (1.66 MB)

Trainable params: 433,898 (1.66 MB)

Non-trainable params: 0 (0.00 B)

Epoch 1/8

107/107 ————— 8s 17ms/step - accuracy: 0.2049 - loss: 2.1071 - val_accuracy: 0.4008 - val_loss: 1.6841

Epoch 2/8

107/107 ————— 1s 11ms/step - accuracy: 0.3758 - loss: 1.7092 - val_accuracy: 0.4993 - val_loss: 1.4561

Epoch 3/8

107/107 ————— 1s 12ms/step - accuracy: 0.5045 - loss: 1.3979 - val_accuracy: 0.5814 - val_loss: 1.1769

Epoch 4/8

107/107 ————— 1s 11ms/step - accuracy: 0.5907 - loss: 1.1596 - val_accuracy: 0.4254 - val_loss: 1.7852

Epoch 5/8

107/107 ————— 1s 11ms/step - accuracy: 0.6294 - loss: 1.0551 - val_accuracy: 0.6922 - val_loss: 0.8964

Epoch 6/8


107/107 ————— 1s 11ms/step - accuracy: 0.6714 - loss: 0.9435 - val_accuracy: 0.6005 - val_loss: 1.0762

Epoch 7/8

107/107 ————— 1s 11ms/step - accuracy: 0.7062 - loss: 0.8705 - val_accuracy: 0.6457 - val_loss: 1.1114

Epoch 8/8


107/107 ————— 2s 18ms/step - accuracy: 0.7147 - loss: 0.8069 - val_accuracy: 0.6347 - val_loss: 1.0787



```
=====
MODEL PERFORMANCE
=====
Test Accuracy: 60.88 %
Test Loss: 1.2648

Classification Report:
```

	precision	recall	f1-score	support
Tomato_Bacterial_spot	1.00	0.34	0.50	80
Tomato_Early_blight	0.43	0.52	0.47	73
Tomato_Late_blight	0.71	0.53	0.61	70
Tomato_Leaf_Mold	0.86	0.59	0.70	85
Tomato_Septoria_leaf_spot	0.72	0.43	0.54	77
Tomato_Spider_mites_Two_spotted_spider_mite	0.54	0.45	0.49	71
Tomato__Target_Spot	0.43	0.62	0.51	73
Tomato__Tomato_YellowLeaf__Curl_Virus	0.88	0.81	0.84	70
Tomato__Tomato_mosaic_virus	0.65	0.90	0.76	52
Tomato_healthy	0.50	0.99	0.66	80
accuracy			0.61	731
macro avg	0.67	0.62	0.61	731
weighted avg	0.68	0.61	0.60	731



CNN Prediction Results

Actual: Tomato_Leaf_Mold
Predicted: Tomato_Septoria_leaf_spot
Confidence: 71.1%



Actual: Tomato_Leaf_Mold
Predicted: Tomato_Leaf_Mold
Confidence: 98.8%



Actual: Tomato_Early_blight
Predicted: Tomato_Early_blight
Confidence: 52.1%



Actual: Tomato_Tomato_mosaic_virus
Predicted: Tomato_Tomato_mosaic_virus
Confidence: 62.6%



Actual: Tomato_Early_blight
Predicted: Tomato_Target_Spot
Confidence: 68.4%



Actual: Tomato_Target_Spot
Predicted: Tomato_Target_Spot
Confidence: 86.0%




Actual: Tomato_healthy
Predicted: Tomato_healthy
Confidence: 98.5%



Actual: Tomato_Target_Spot
Predicted: Tomato_Target_Spot
Confidence: 89.4%



Actual: Tomato_Tomato_mosaic_virus
Predicted: Tomato_Tomato_mosaic_virus
Confidence: 94.3%



Actual: Tomato_healthy
Predicted: Tomato_healthy
Confidence: 98.5%



Actual: Tomato_Target_Spot
Predicted: Tomato_Target_Spot
Confidence: 89.4%



Actual: Tomato_Tomato_mosaic_virus
Predicted: Tomato_Tomato_mosaic_virus
Confidence: 94.3%



Model saved successfully.